Data Storytelling Final Project

Mitchell Beckner

3/28/2020

Author

Mitchell Beckner

Links

Final Project Markdown on GitHub:

Presentation Deck .pptx on GitHub:

Data Cleaning and EDA in R: https://github.com/mitchb63/MovieLens_Analysis

Summary

This project was created as part of Udacity's Data Visualization Nanodegree program. The data used is from a Kaggle dataset that contains metadata for 45,000 movies listed in the MovieLens Dataset. The dataset consists of movies released on or before July 2017. Data points include cast, crew, plot keywords, budget, revenue, posters, release dates, languages, production companies, countries, TMDB vote counts and vote averages. The dataset also has 26 million ratings from 270,000 users for all 45,000 movies (Banik, 2017).

The objective of this project was to use the movies dataset and conduct the EDA necessary to understand the dataset as a whole. The output will be to discover if the dataset is balanced, if there are anomalies in the dataset that affect the applicability of the recommendation, and the final presentation itself that will be used for a mock recommendation to a management team (Udacity, 2020).

The problem statement for the project was defined as:

What are the key components that differentiate the top 20% of films from the remainder when considering profit margin?

Three hypotheses were constructed as the foundation of the study: * Story factors influence the film's profit margin * Production factors influence the profit margin * Public opinion factors affect the profit margin

The key findings of the analysis were that Story factors and budget appear to have the greatest influence on the profit margin of a film and the final recommendations were to focus on the production and promotion of films that are part of an ongoing collection and/or films with budgets less than approximately $6 million. These movies were found to be more likely to have profit margins in the top 20% of those examined. The analytical process used, as well as the limitations and biases in the dataset, are described in the sections that follow.

Analysis Process

The primary dataset was downloaded from the Kaggle website (Banik, 2017). It consisted of 7 data files. Additional data was downloaded from the IMDB website that consisted of lists of the highest rated actors, producers and directors (nims-1975, 2015)(Smmsadrnezh, 2018). This information was used as reference to create binary variables for 'use of top-rated cast' and 'use of top-rated crew' for each film.

Some preliminary data preparation and cleaning was done in Excel before importing the data into R for further processing. Many of the "0" values were replaced with NA's in order to accurately account for, and later impoute the missing data. Duplicate rows were also discovered and removed. A summary of the distributions of the variables in the dataset is shown in Figure 1.

```
> summary(df_movies_init)
       id             budget          original_language   popularity       release_date          revenue             runtime
 2      :    1   Min.   :        1   en     :32251     Min.   :  0.000   Min.   :1874-12-09   Min.   :          1   Min.   :   1.0
 3      :    1   1st Qu.:  2000000   fr     : 2436     1st Qu.:  0.386   1st Qu.:1978-10-06   1st Qu.:   2400000   1st Qu.:  86.0
 5      :    1   Median :  8000000   it     : 1529     Median :  1.127   Median :2001-08-30   Median :  16829545   Median :  95.0
 6      :    1   Mean   : 21614181   ja     : 1347     Mean   :  2.921   Mean   :1992-05-16   Mean   :  68489031   Mean   :  97.5
 11     :    1   3rd Qu.: 25000000   de     : 1079     3rd Qu.:  3.678   3rd Qu.:2010-12-17   3rd Qu.:  67308282   3rd Qu.: 107.0
 (Other):45428   Max.   :380000000   (other): 6781     Max.   :547.488   Max.   :2020-12-16   Max.   :2068223624   Max.   :1256.0
 NA's   :    1   NA's   :36554       NA's   :   11     NA's   :4         NA's   :88           NA's   :38037        NA's   :1819
           status                title        vote_average     vote_count    release_month    release_year
 Canceled       :    2   Cinderella        :   11   Min.   : 0.50   Min.   :    0   Jan    : 5909   2014   : 1973
 In Production  :   20   Alice in Wonderland :    9   1st Qu.: 5.30   1st Qu.:    3   Sep    : 4834   2015   : 1904
 Planned        :   15   Hamlet            :    9   Median : 6.10   Median :   10   Oct    : 4613   2013   : 1887
 Post Production:   98   Beauty and the Beast:    8   Mean   : 6.01   Mean   :  110   Dec    : 3781   2012   : 1721
 Released       :44985   Les Misérables    :    8   3rd Qu.: 6.90   3rd Qu.:   34   Nov    : 3661   2011   : 1666
 Rumored        :  229   (Other)           :45385   Max.   :10.00   Max.   :14075   (Other):22548   (Other):36195
 NA's           :   85   NA's              :    4   NA's   :2999    NA's   :4       NA's   :   88   NA's   :   88
>
```

Figure 1: Raw Data Distribution Summary

Next the dataset was filtered to include only films that had been released to the public and those with budgets less than $59,500,000. This eliminated the statistical outliers found in the budget variable. New variables for *profit* and *profit_margin* were created and then used to create a *profit_group* categorical variable that placed each movie into either the top 20% or bottom 80% based on profit margin. Finally, R code was written to process the separate files for ratings, keywords, cast and crew data, country and language data, genre, and movie collection information. The full code is available at the link posted above.

Various R packages were used to explore the missingness present within the data. Figure 2 shows a plot of the pattern of missing data and Figure 3 shows the patterns of missingness. It should be noted that while the original dataset contained information on over 45,000 movies, removing duplicates and filtering for 'Released" movies with a stated budget greater than 0 resulted in only 7925 films. These films were used as the basis for the remaining investigation.

The 'MICE" package was then used to impute the missing values in the dataset (Noghrehchi, 2015). Five imputed datasets were created using 50 iterations of the chained equation process. The results were examined for convergence and found to be satisfactory. A single imputed dataset was then randomly selected and used for the remainder of the analysis.

Following the imputation of missing values, the variables for profit and profit_margin were re-generated and again used to create the profit_group variable that placed each movie into either the top 20% or bottom 80% based on profit margin. A summary of the variable distributions for the imputed dataset are shown in Figure 4. This data was then standardized and correlation analysis was run. An excerpt of the results of this analysis is shown in Figure 5 with the corresponding p-values shown in Figure 6. These results showed significant correlations betw.

Following the correlation analysis, additional EDA was performed in R to confirm the results and lay the groundwork for the final visualizations. Final visualizations were then produced in Tableau and R with some additional polishing done in Adobe Illustrator. The final presentation was then assembled based on the PowerPoint ghost deck created in a prior project.

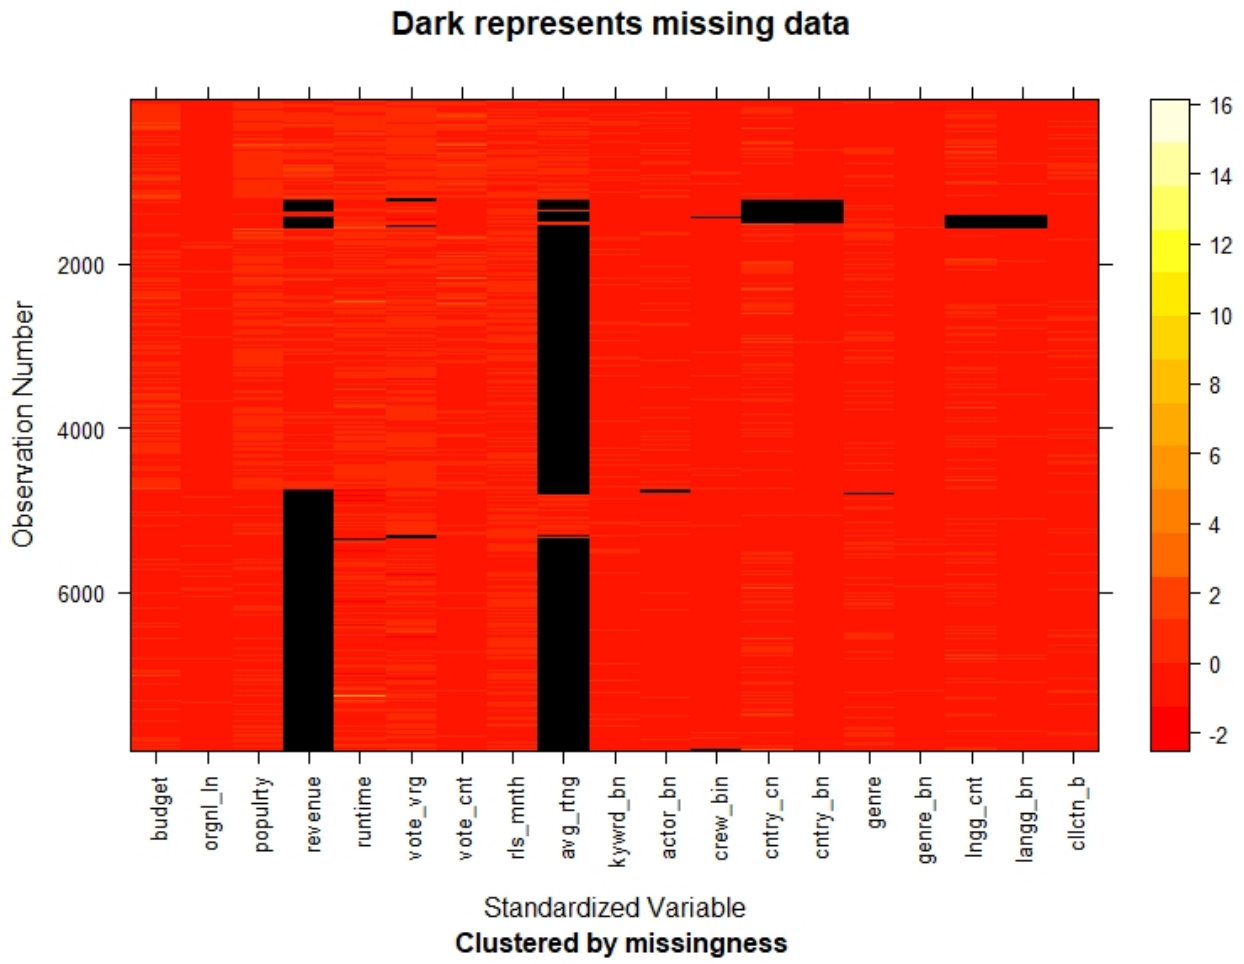

Figure 2: Missing Data as Shown in MI

```
> show(missing_bins)
Object of class missing_data.frame with 7925 observations on 20 variables

There are 95 missing data patterns

Append '@patterns' to this missing_data.frame to access the corresponding pattern for every observation or perhaps use table()

                                type missing method  model
id                          irrelevant       0   <NA>   <NA>
budget                      continuous       0   <NA>   <NA>
original_language unordered-categorical      0   <NA>   <NA>
popularity                  continuous       0   <NA>   <NA>
revenue                     continuous    3435    ppd linear
runtime                     continuous      66    ppd linear
vote_average                continuous      90    ppd linear
vote_count                  continuous       0   <NA>   <NA>
release_month     unordered-categorical      4    ppd mlogit
avg_rating                  continuous    6151    ppd linear
keyword_bin                     binary       0   <NA>   <NA>
actor_bin                       binary      77    ppd  logit
crew_bin                        binary      37    ppd  logit
country_count               continuous     286    ppd linear
country_bin                     binary     286    ppd  logit
genre             unordered-categorical      57    ppd mlogit
genre_bin                       binary       0   <NA>   <NA>
language_count              continuous     165    ppd linear
language_bin                    binary     165    ppd  logit
collection_bin                  binary       0   <NA>   <NA>
```

Figure 3: Missing Data Summary

```
> summary(df_imp_1)
       id              budget          original_language   popularity        revenue             runtime      vote_average     vote_count
 5      :   1   Min.   :        1   en     :6497      Min.   :  0.00   Min.   :        1   Min.   :  4   Min.   : 0.50   Min.   :    0
 11     :   1   1st Qu.: 1525000   fr     : 234      1st Qu.:  1.84   1st Qu.:  1477676   1st Qu.: 91   1st Qu.: 5.40   1st Qu.:   17
 13     :   1   Median : 6229577   ru     : 136      Median :  5.62   Median :  9366227   Median :100   Median : 6.20   Median :   71
 14     :   1   Mean   :12048609   hi     : 130      Mean   :  6.33   Mean   : 32164871   Mean   :105   Mean   : 6.06   Mean   :  298
 15     :   1   3rd Qu.:19000000   es     : 113      3rd Qu.:  9.19   3rd Qu.: 33697647   3rd Qu.:114   3rd Qu.: 6.80   3rd Qu.:  285
 16     :   1   Max.   :59000000   de     :  94      Max.   :228.03   Max.   :792965326   Max.   :523   Max.   :10.00   Max.   :11444
 (Other):7919                      (Other): 721
 release_month    avg_rating    keyword_bin actor_bin crew_bin country_count  country_bin       genre       genre_bin language_count
 Sep     : 958   Min.   :0.50   0:7139      0:6920    0:7787   Min.   :1.00   0:7825      Drama    :2092   0:7887    Min.   :1.0
 Oct     : 830   1st Qu.:2.83   1: 786      1:1005    1: 138   1st Qu.:1.00   1: 100      Comedy   :1590   1:  38    1st Qu.:1.0
 Dec     : 714   Median :3.24                                  Median :1.00               Action   :1174             Median :1.0
 Jan     : 705   Mean   :3.15                                  Mean   :1.34               Horror   : 710             Mean   :1.4
 Aug     : 674   3rd Qu.:3.56                                  3rd Qu.:1.00               Crime    : 383             3rd Qu.:2.0
 Apr     : 618   Max.   :4.50                                  Max.   :5.00               Adventure: 375             Max.   :5.0
 (Other):3426                                                                             (Other)  :1601
 language_bin collection_bin     profit          profit_margin       profit_group
 0:7587       0:6689         Min.   :-53223667   Min.   :-1999999   Bottom_80:6340
 1: 338       1:1236         1st Qu.: -1862635   1st Qu.:      -1   Top_20   :1585
                             Median :  1923000   Median :       0
                             Mean   : 20116262   Mean   :   -4719
                             3rd Qu.: 18907422   3rd Qu.:       1
                             Max.   :782465326   Max.   :       1
```

Figure 4: Imputed Data Distribution Summary

```
> corr <- round(cor(df_imp1_std, use="pairwise.complete.obs"), 6)
> corr
               profit_group profit_margin    budget popularity   revenue   runti
profit_group       1.000000      0.037199 -0.238984   0.045164  0.278838 -0.1010
profit_margin      0.037199      1.000000  0.044505   0.042649  0.038395  0.0133
budget            -0.238984      0.044505  1.000000   0.352454  0.482830  0.2048
popularity         0.045164      0.042649  0.352454   1.000000  0.422242  0.1111
revenue            0.278838      0.038395  0.482830   0.422242  1.000000  0.1409
runtime           -0.101066      0.013334  0.204837   0.111162  0.140948  1.0000
vote_average       0.101696      0.030254  0.041441   0.233059  0.184336  0.2569
vote_count         0.150940      0.031343  0.354916   0.592744  0.650708  0.1130
avg_rating         0.022407     -0.011500 -0.023558   0.016004  0.033491  0.0344
keyword_bin        0.057835      0.002307  0.027240   0.080756  0.088348  0.0060
actor_bin         -0.048348      0.026336  0.277909   0.182051  0.171968  0.1397
crew_bin           0.017846      0.009899  0.077829   0.093790  0.134254  0.0731
country_count     -0.092008      0.003703  0.114902   0.069720  0.004748  0.0940
country_bin        0.005652      0.002186 -0.038108  -0.030632 -0.029949  0.0140
genre_bin         -0.007307     -0.000057 -0.037219  -0.031921 -0.019363 -0.0032
language_count    -0.051013     -0.003741  0.116051   0.094269  0.065350  0.1637
language_bin       0.005308     -0.016623  0.003511   0.013190  0.012173  0.0702
collection_bin     0.120681      0.023916  0.089118   0.170178  0.255433 -0.0508
profit             0.362190      0.031314  0.288230   0.377213  0.977716  0.1050
```

Figure 5: Data Correlation Summary

```
> p.mat <- round(cor_pmat(df_imp1_std), 6)
> p.mat
              profit_group profit_margin   budget popularity  revenue  runt
profit_group      0.000000      0.000926 0.000000   0.000058 0.000000 0.000
profit_margin     0.000926      0.000000 0.000074   0.000146 0.000629 0.235
budget            0.000000      0.000074 0.000000   0.000000 0.000000 0.000
popularity        0.000058      0.000146 0.000000   0.000000 0.000000 0.000
revenue           0.000000      0.000629 0.000000   0.000000 0.000000 0.000
runtime           0.000000      0.235288 0.000000   0.000000 0.000000 0.000
vote_average      0.000000      0.007071 0.000224   0.000000 0.000000 0.000
vote_count        0.000000      0.005263 0.000000   0.000000 0.000000 0.000
avg_rating        0.046082      0.305994 0.035983   0.154271 0.002865 0.002
keyword_bin       0.000000      0.837279 0.015306   0.000000 0.000000 0.588
actor_bin         0.000017      0.019051 0.000000   0.000000 0.000000 0.000
crew_bin          0.112152      0.378264 0.000000   0.000000 0.000000 0.000
country_count     0.000000      0.741735 0.000000   0.000000 0.672550 0.000
country_bin       0.614887      0.845717 0.000691   0.006388 0.007669 0.209
genre_bin         0.515464      0.995930 0.000920   0.004484 0.084773 0.774
language_count    0.000006      0.739171 0.000000   0.000000 0.000000 0.000
language_bin      0.636603      0.138957 0.754675   0.240353 0.278564 0.000
collection_bin    0.000000      0.033252 0.000000   0.000000 0.000000 0.000
profit            0.000000      0.005305 0.000000   0.000000 0.000000 0.000
```

Figure 6: p-Values for Correlated Data

Limitations and Biases

###Data Collection Phase

Many problems were found in the provided dataset. As noted above, the data contained duplicate rows, and a large amount of missing data. In addition, some the data that was present was obviously erroneous and/or suspect. For example, some movies had negative values listed as their budget. Without more information on the methods of data collection and their sources, it is impossible to determine the reasons for these errors and omissions, therefore, assumptions that were made regarding the data being missing completely at random are essentially indefensible. It is also likely that a great deal of response bias is present in the popularity, vote average and rating data since this information was presumably collected at various lengths of time after a film was released.

Data Processing Phase

Again, as noted above, the MICE package was used to impute missing data in R. While every attempt was made to select the best available algorithm for each variable, and results were visually examined to ensure that the imputed values were 'logically reasonable, there is the potential for a great deal of bias being introduced at this point in the process. In addition, in an attempt to use the cast and crew data that was provided, additional lists of "top-rated" actors, producers, and directors were obtained from the IMDB website. Each movies cast and crew was checked against these lists and movies were flagged if they used actors or crew members on these outside lists. However, there is no standard criteria that was used to establish the "top-rated" status of these people so again, bias was likely introduced at this point in the analysis.

Insight Phase

While every attempt was made to verify results and recommendations by multiple methods including correlation analysis, creation of visualizations, and careful analysis of potential biases, it is possible that confirmation

bias is present to a degree in the findings. In addition, were a model to be created in order to predict profit margin based on the variables determined to be significant, it is possible that the model would be over or under fitted given the relatively small percentage of the dataset that began with complete information.

References

Anonymous. (1970, January 1). McKinsey Presentations - How to Apply Ghost (aka Shell and Skeleton) Decks and Pages. Retrieved March 7, 2020, from http://workingwithmckinsey.blogspot.com/2013/07/McKinsey-presentations-ghost-decks.html

Banik, R. (2017, November 10). The Movies Dataset. Retrieved March 1, 2020, from https://www.kaggle.com/rounakbanik/the-movies-dataset#movies_metadata.csv

nims-1975. (2015, July 2). Top 12 Best Film Producers/Directors. Retrieved from https://www.imdb.com/list/ls051116454/

Noghrehchi, F. (2015, May 29). Retrieved from https://web.maths.unsw.edu.au/~dwarton/missingDataLab.html

Smmsadrnezh. (2018, March 5). Best and Worst Actors Rating (Sorted). Retrieved from https://www.imdb.com/list/ls070682726/

Udacity. (2020). Build a Data Story Final Project. Retrieved from https://classroom.udacity.com/nanodegrees/nd197/parts/c26016f6-de16-42b7-98f8-9de7374f8247/modules/2c1ae2e0-f442-4268-9931-c8b2d0c07ea2/lessons/060138e7-74ee-4391-8ced-2f1fbcca002c/concepts/1e50f104-b96f-462a-a408-a155ea4a2467